

# psiflow

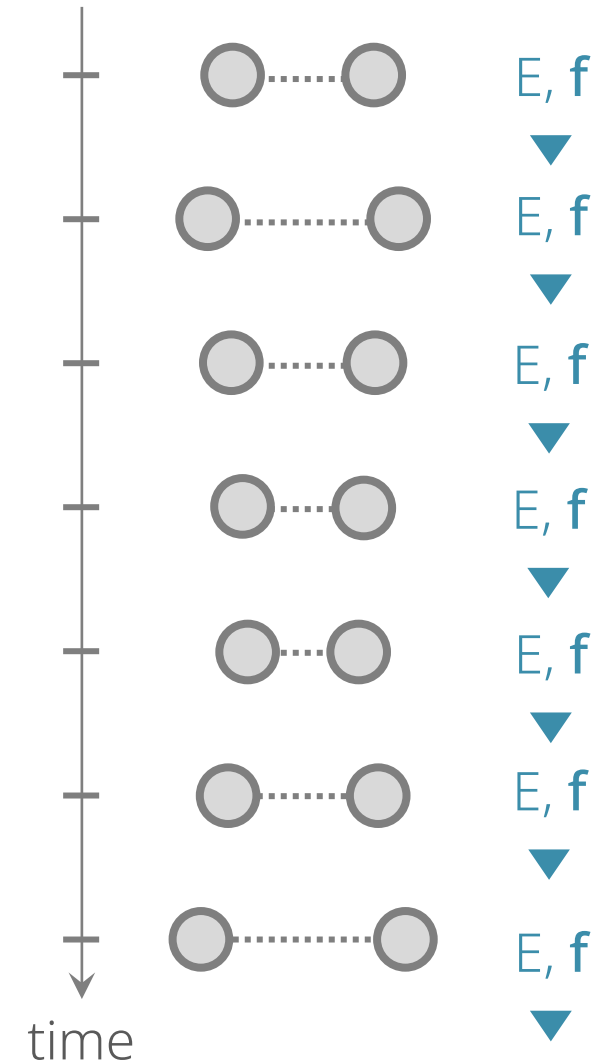
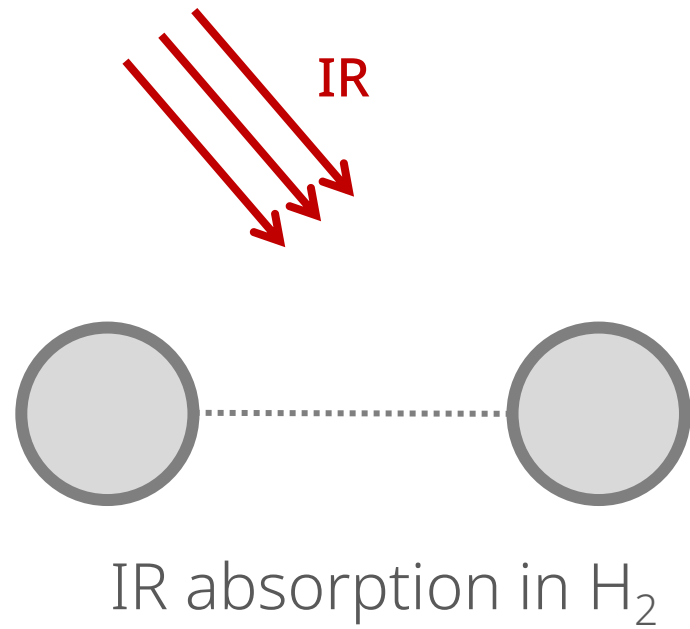
modular and scalable online learning  
for atomic simulations

Sander Vandenhoute

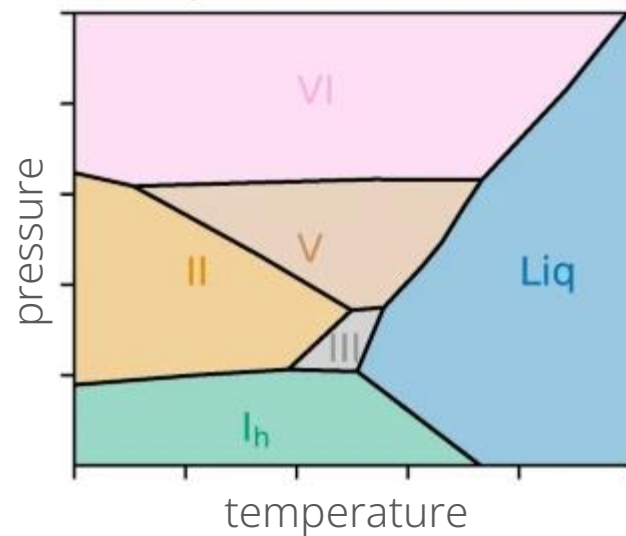
Prof. Veronique Van Speybroeck

Ghent University (BE)

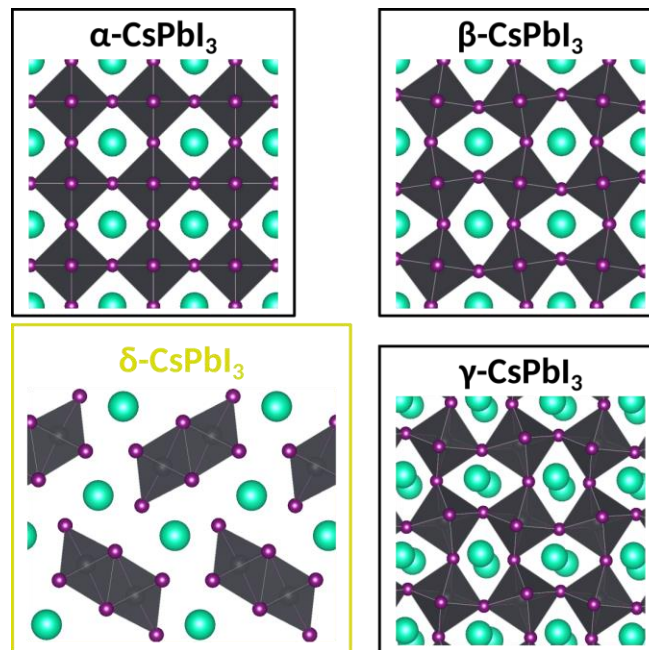
# physical property



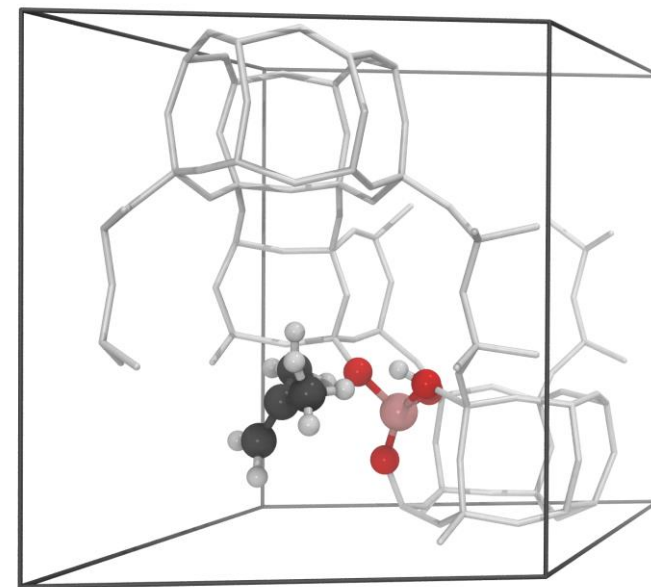
phase diagrams



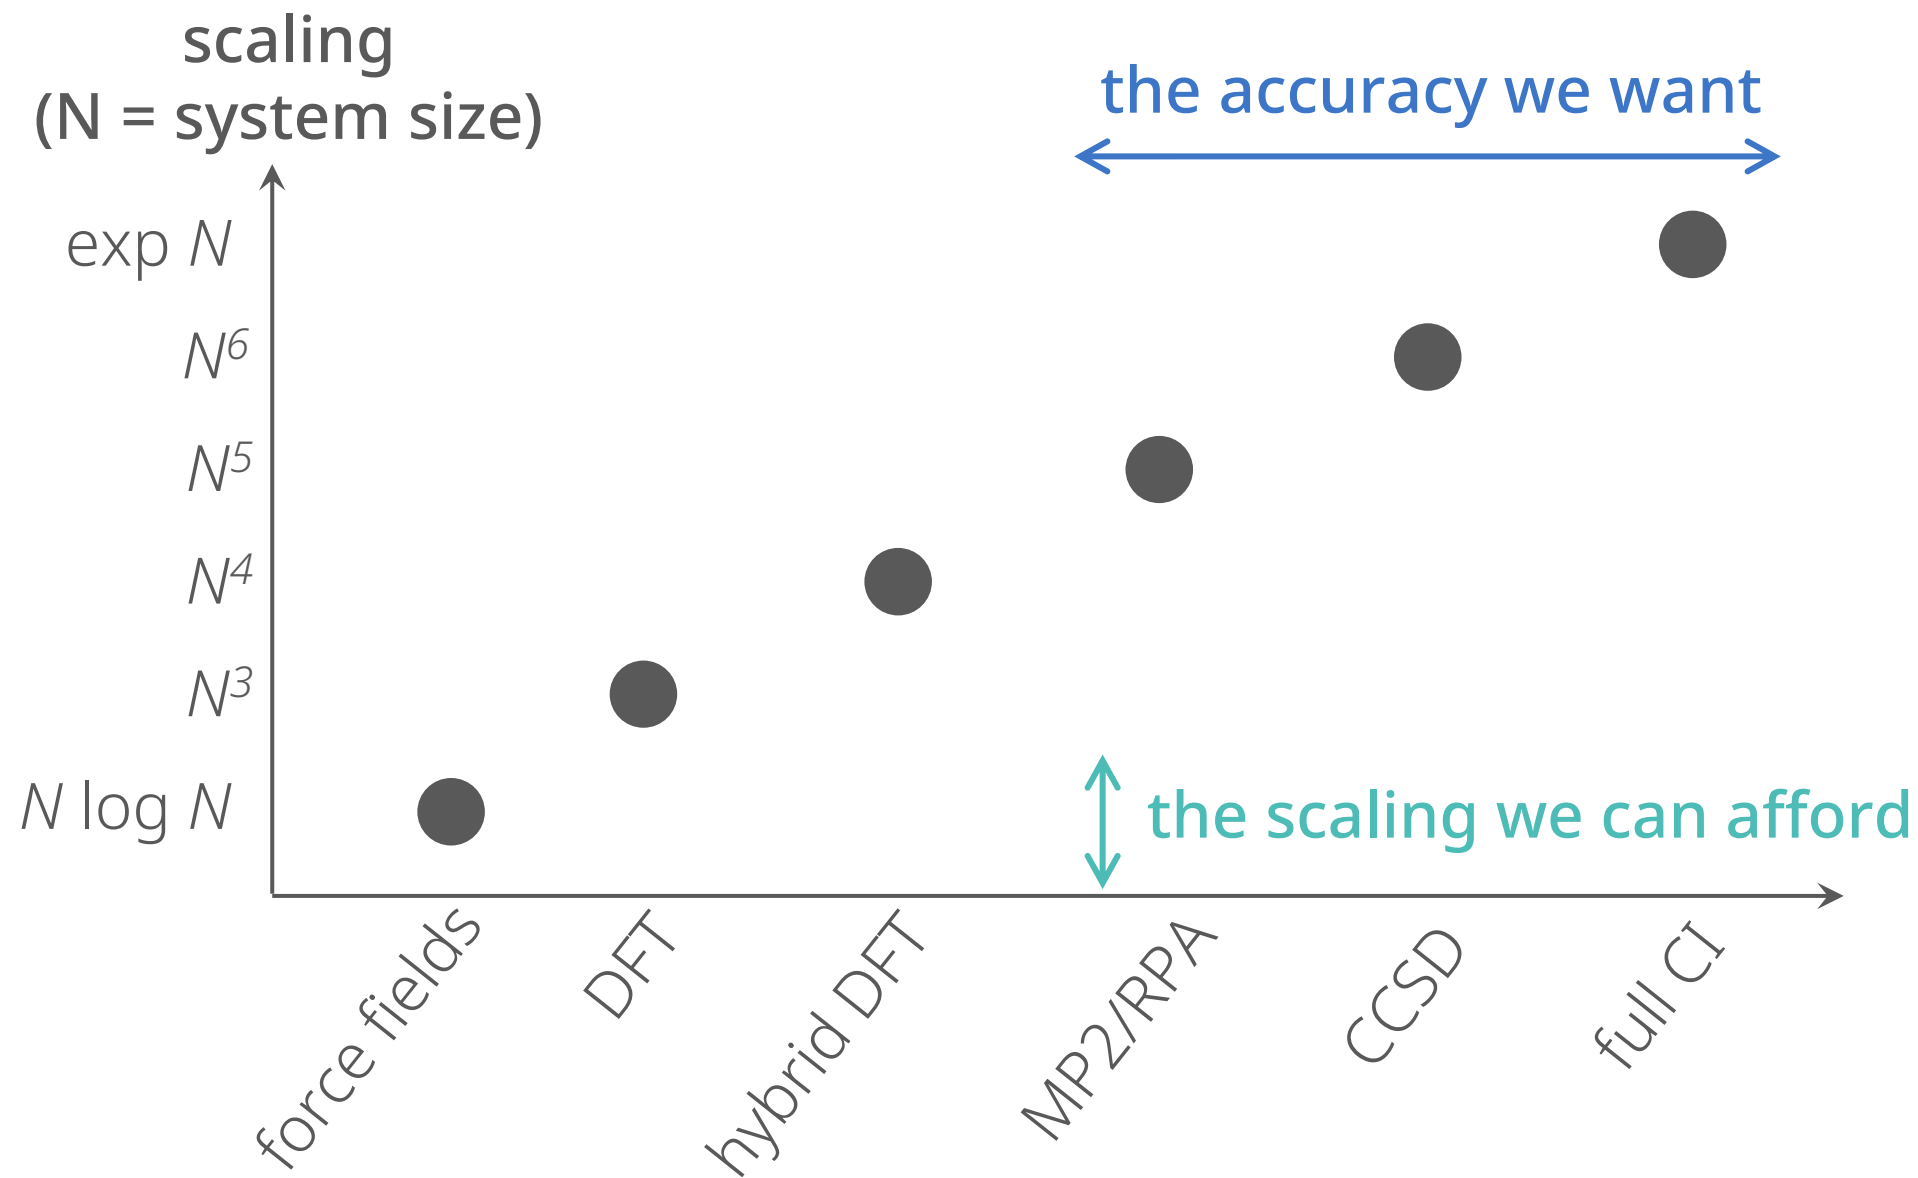
thermodynamic stability



chemical reactivity

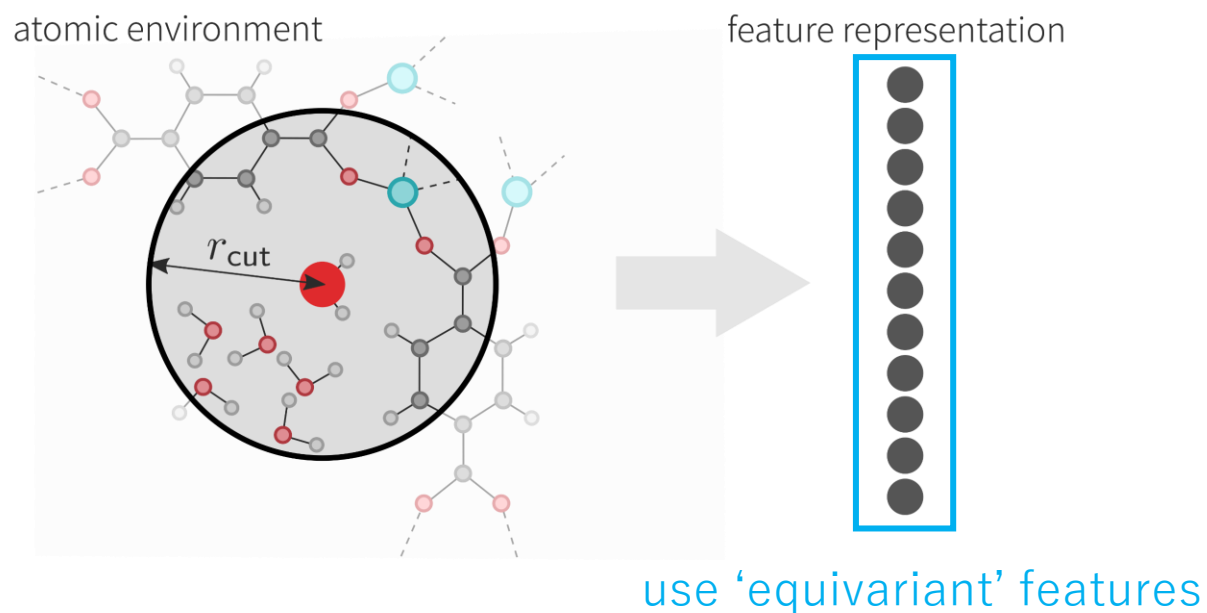


~~hundreds~~  
~~thousands~~  
millions/billions of E/f evaluations

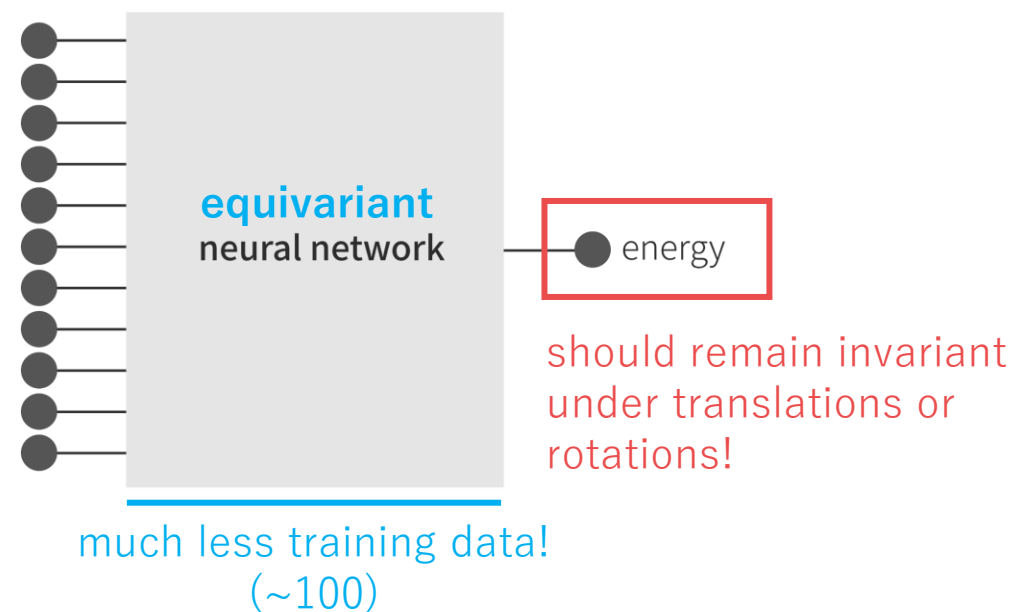


# Use equivariant neural networks to learn E/f

## STEP 1



## STEP 2



```
while error(model) is high:
```

```
    data = generate_atomic_data(model)           # CPU/GPU
```

```
    data = evaluate_DFT_energy_forces(data)      # CPU
```

```
    model.train(data)                             # GPU
```

```
evaluate_exact_groundstate_energy(data)         # large-memory CPU
```

```
model.train(data)                               # GPU
```

```
compute_free_energies(model)                   # CPU/GPU
```



average computational chemist is not very computational

→ hide Parsl API as much as possible

enforce “write once, run anywhere”

→ separate configuration from high-level workflow definition

# WHAT?

high-level workflow definition  
using abstractions:

Model  
Dataset  
Walker  
Learning  
...

---

run.py

# HOW?

run CP2K on N cores

train model for M minutes

Docker/Apptainer URIs

MPI/OpenMP

---

lumi.py  
frontier.py

# WHERE?

Parsl providers!

```
python run.py                # runs locally using ThreadPoolExecutors
python run.py --psiflow-config lumi.py    # on LUMI
```



```
import psiflow

from psiflow.data import Dataset      # wraps around File/DataFuture
from psiflow.models import MACEModel # wraps around File/DataFuture

psiflow.load() # build and load Parsl config; 'compile' apps

train, valid = Dataset.load('dft.xyz').split(0.9)

model = MACEModel() # contains DataFuture of untrained model
model.train(train, valid) # contains DataFuture of trained model

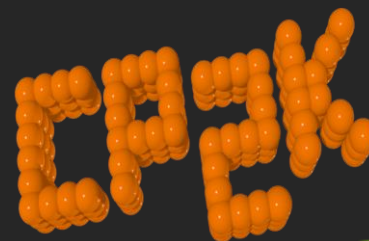
errors = Dataset.get_errors(valid, model.evaluate(valid)) # Future
errors.result() # NumPy array of validation errors on energy/force
```



OpenMM



PyTorch



NWChem



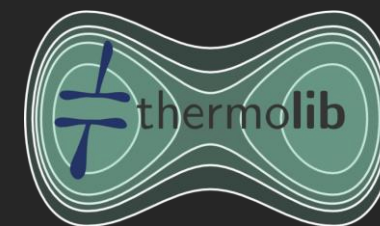
Weights & Biases



MACE



PLUMED



bundled in a Docker/Apptainer image → **ContainerizedLauncher**

**Big workflows create huge amounts of files**

→ automatic tarring? Or even archived by default?

**Big workflows require a lot of memory**

→ more extensive use of virtual, on-disk memory?

**Debugging Parsl workflows is nontrivial**

→ lazy failure not always best option?





Veronique Van Speybroeck

Massimo Bocus  
Tom Braeckeveld  
Pieter Dobbelaere  
& others



**EuroHPC**  
Joint Undertaking

**L U M I**



**Research Foundation  
Flanders**  
Opening new horizons

**VLAAMS  
SUPERCOMPUTER  
CENTRUM**



**Vlaanderen**  
is supercomputing



**European  
Research  
Council**